

# Reinforcement Learning III

Lecture 21

# Markov Models

States are  
**Fully Observable**

States are  
**Partially Observable**

**Autonomous**  
(no actions;  
make predictions)

Markov Chain,  
Markov Reward Process

Hidden Markov Model  
(HMM)

**Controlled**  
(can take actions)

Markov Decision  
Process (MDP)

Partially Observable  
Markov Decision  
Process (POMDP)

## Applications

HMMs: time series ML, e.g. speech + handwriting recognition, bioinformatics

MDPs: used extensively for reinforcement learning

# Building blocks for the full RL problem

1	Markov Chain	{state space $S$ , transition probabilities $P$ }
2	Markov Reward Process (MRP)	{ $S$ , $P$ , + rewards $R$ , discount rate $\gamma$ } adds rewards (and values)
3	Markov Decision Process (MDP)	{ $S$ , $P$ , $R$ , $\gamma$ , + actions $A$ } adds decisions (i.e. the ability to control)

**MDPs form the basis for most reinforcement learning environments**

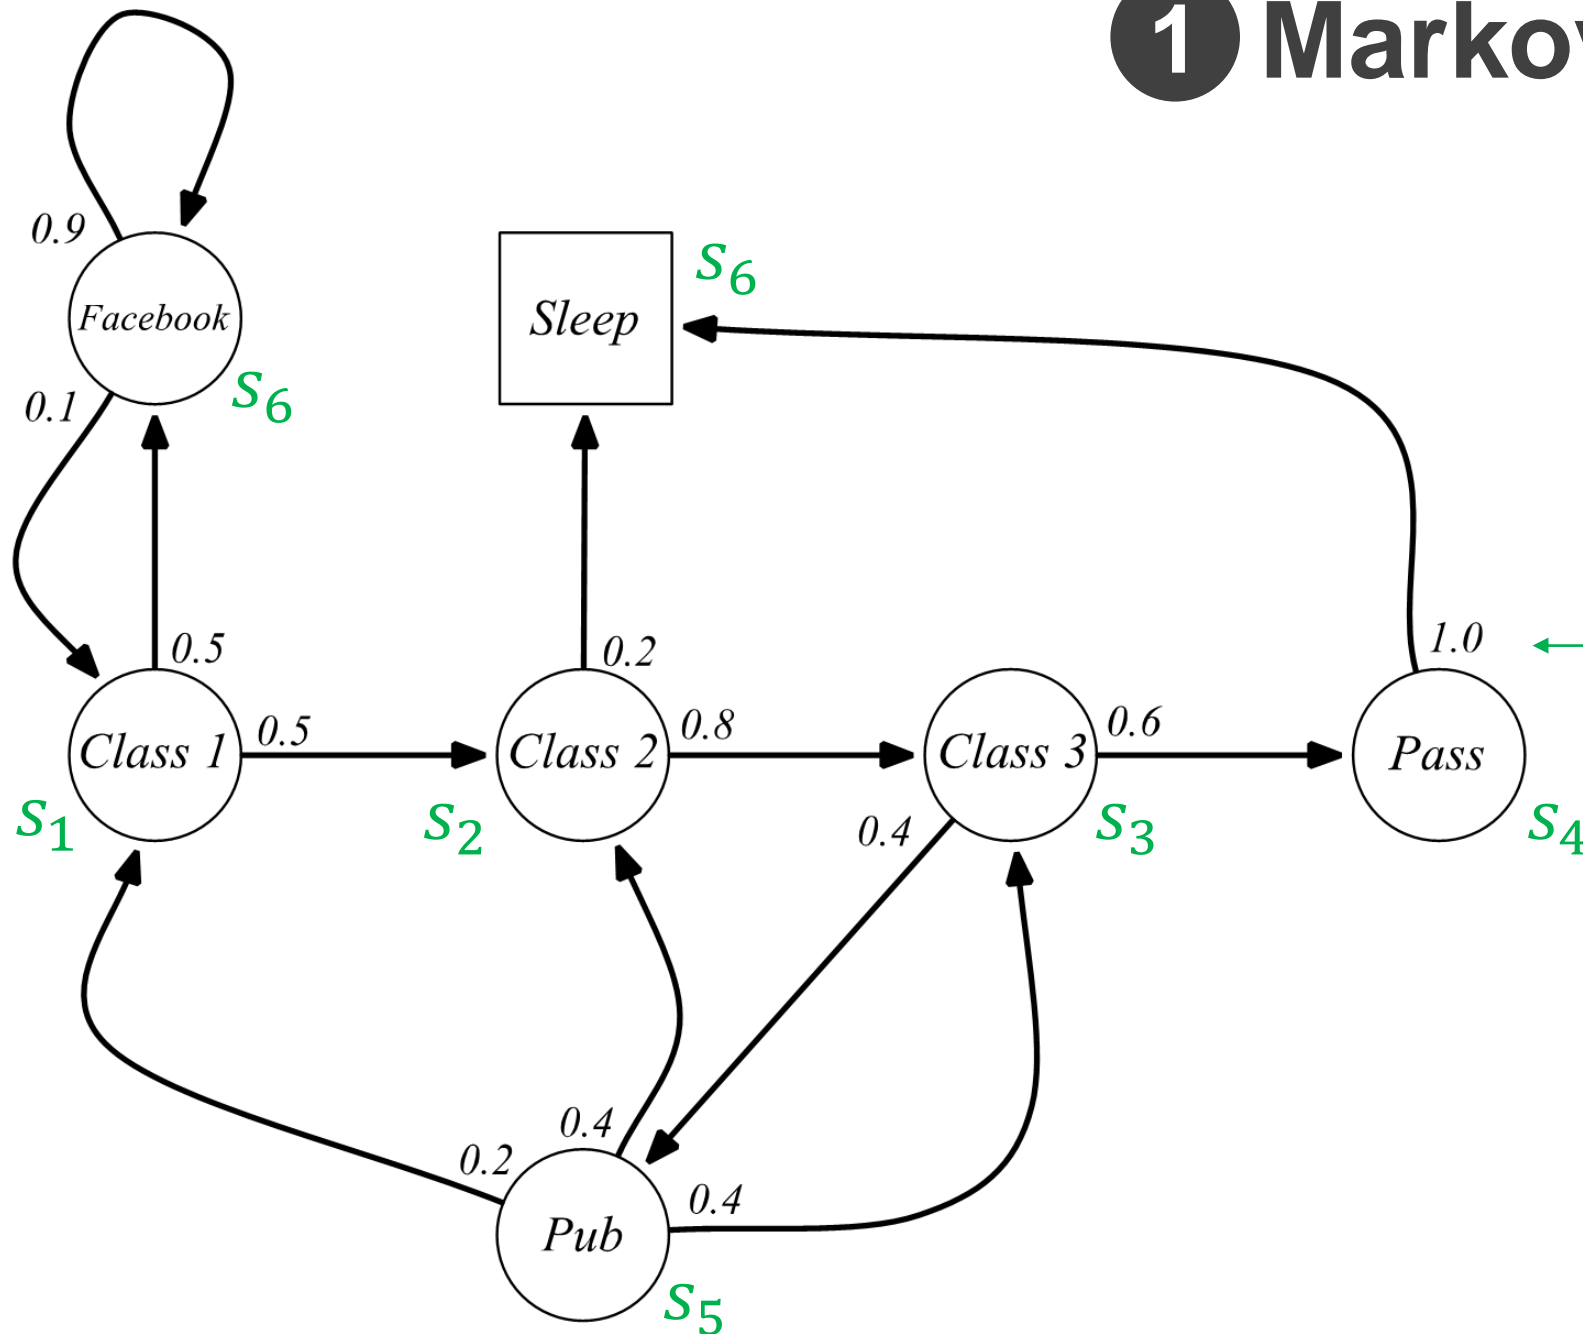
Adapted from David Silver, 2015

# 1 Markov Chain Example

**Components:**

State space  $S$ ,

Transition probabilities  $P$



$$\leftarrow P_{46} = P_{ss'}$$

**Sample Episodes:**

C1,C2,Sleep

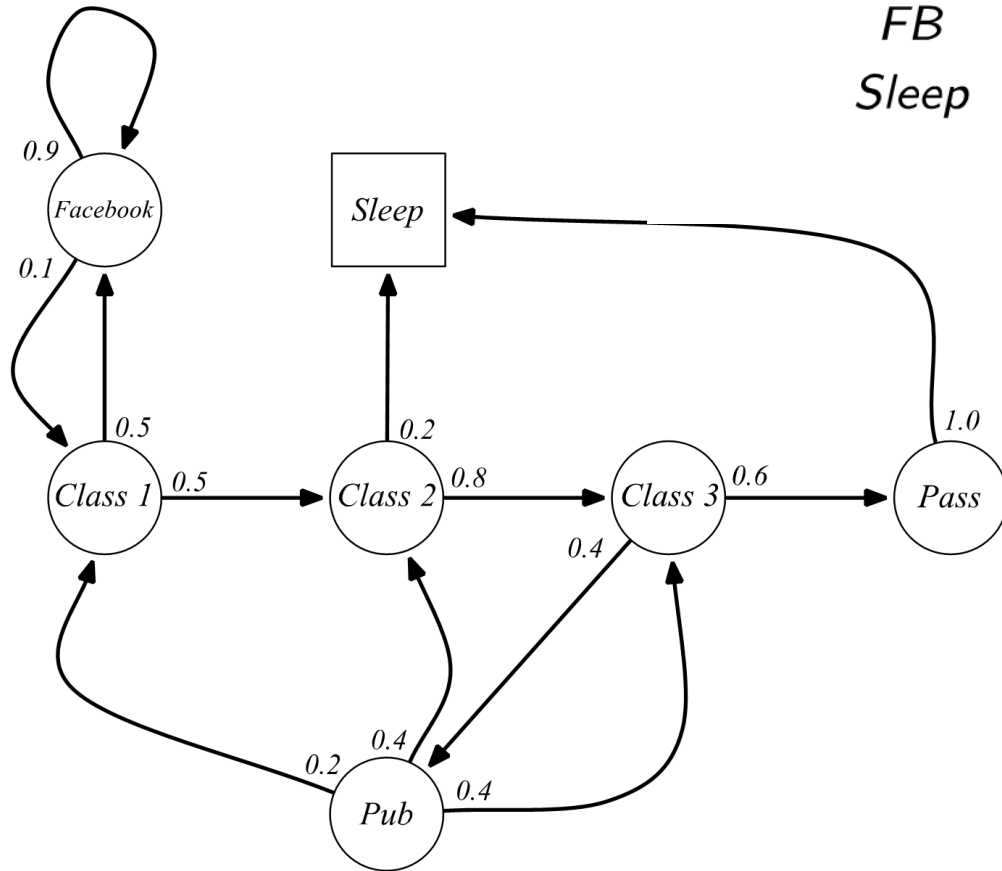
C1,FB,FB,FB,C1,C2,C3,Pass,Sleep

Example from David Silver, UCL, 2015

# Markov Chain

$$\mathcal{P} = \begin{matrix} & \begin{matrix} C1 \\ C2 \\ C3 \\ Pass \\ Pub \\ FB \\ Sleep \end{matrix} \\ \begin{matrix} C1 \\ C2 \\ C3 \\ Pass \\ Pub \\ FB \\ Sleep \end{matrix} & \begin{bmatrix} & & & & & & \\ & 0.5 & & & & 0.5 & \\ & & 0.8 & & & & 0.2 \\ & & & 0.6 & 0.4 & & \\ 0.2 & 0.4 & 0.4 & & & & \\ 0.1 & & & & & 0.9 & \\ & & & & & & 1 \end{bmatrix} \end{matrix}$$

State transition probability matrix,  $P_{ss'}$



Example from David Silver, UCL, 2015

## 2 Markov Reward Process

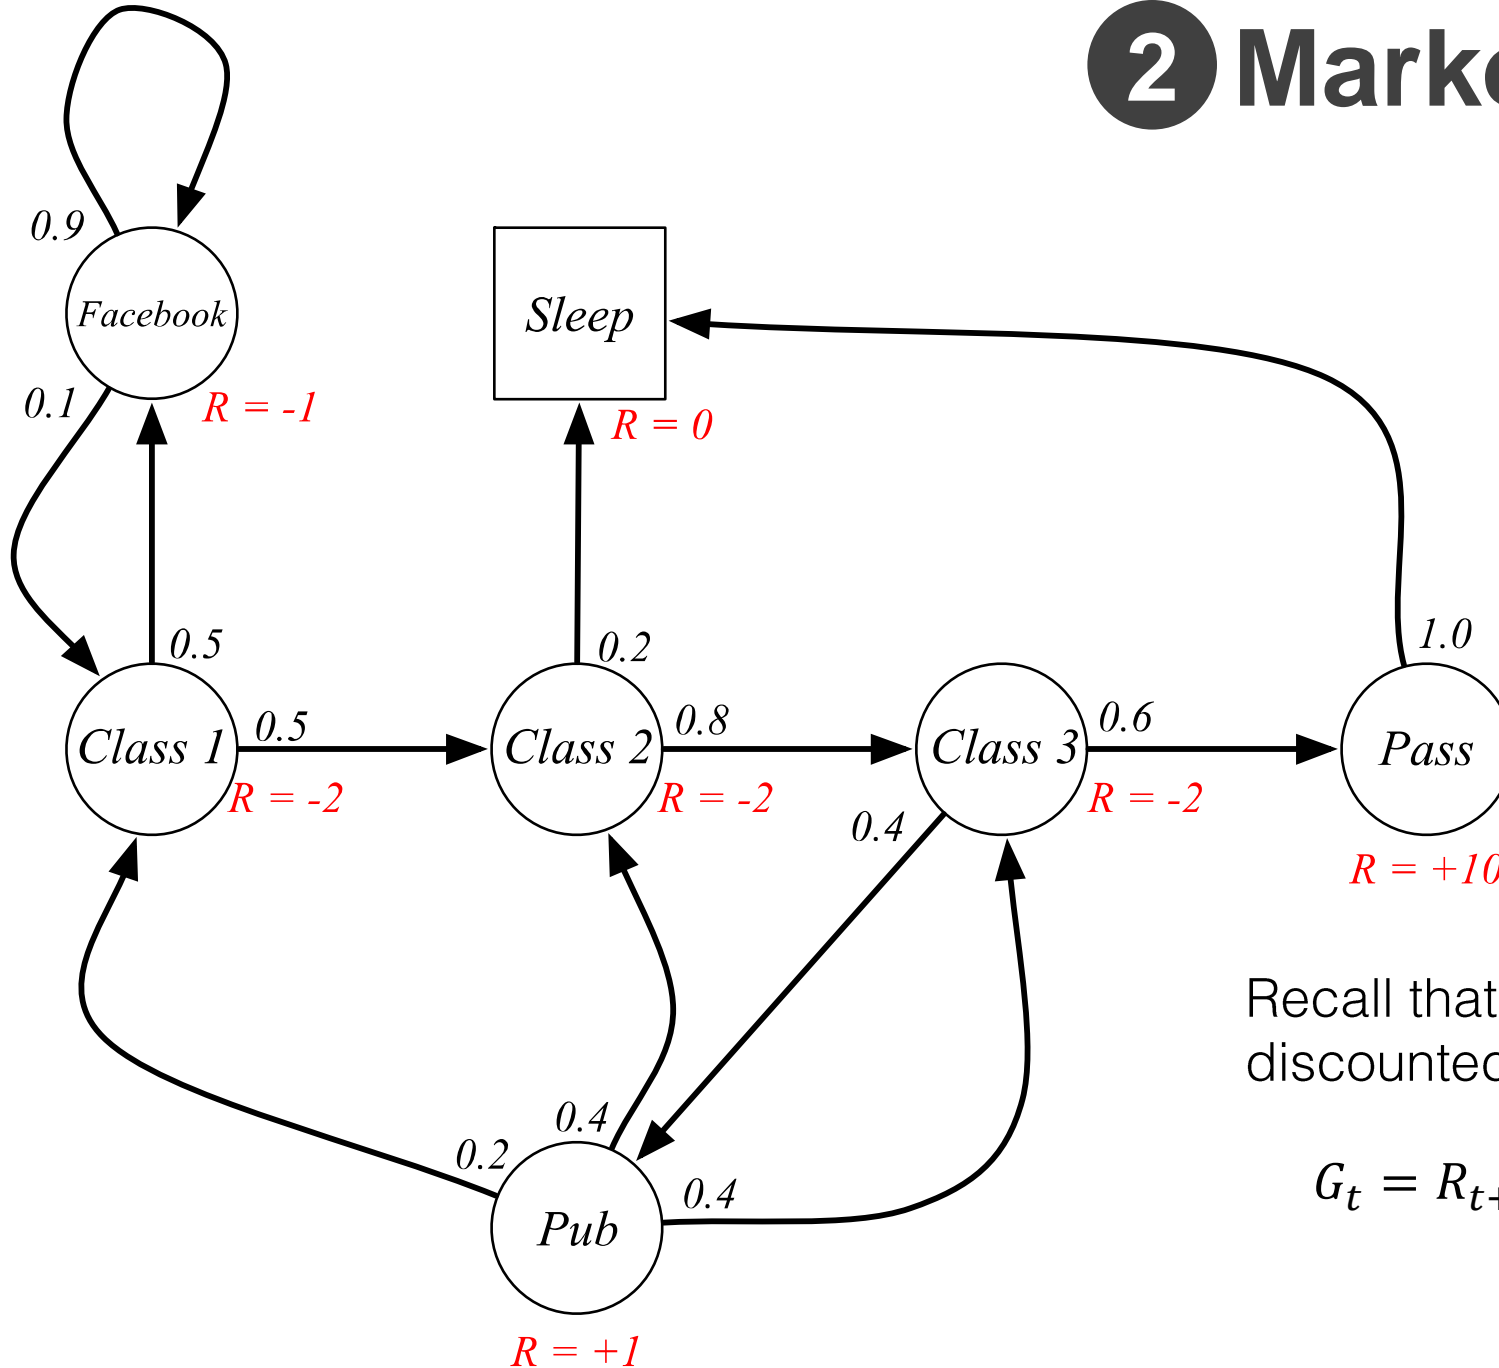
### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



Recall that returns, let's call  $G_t$ , are the total discounted rewards from time  $t$ :

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots = \sum_{k=0}^{\infty} \gamma^k R_{t+k+1}$$

Example from David Silver, UCL, 2015

## 2 Markov Reward Process

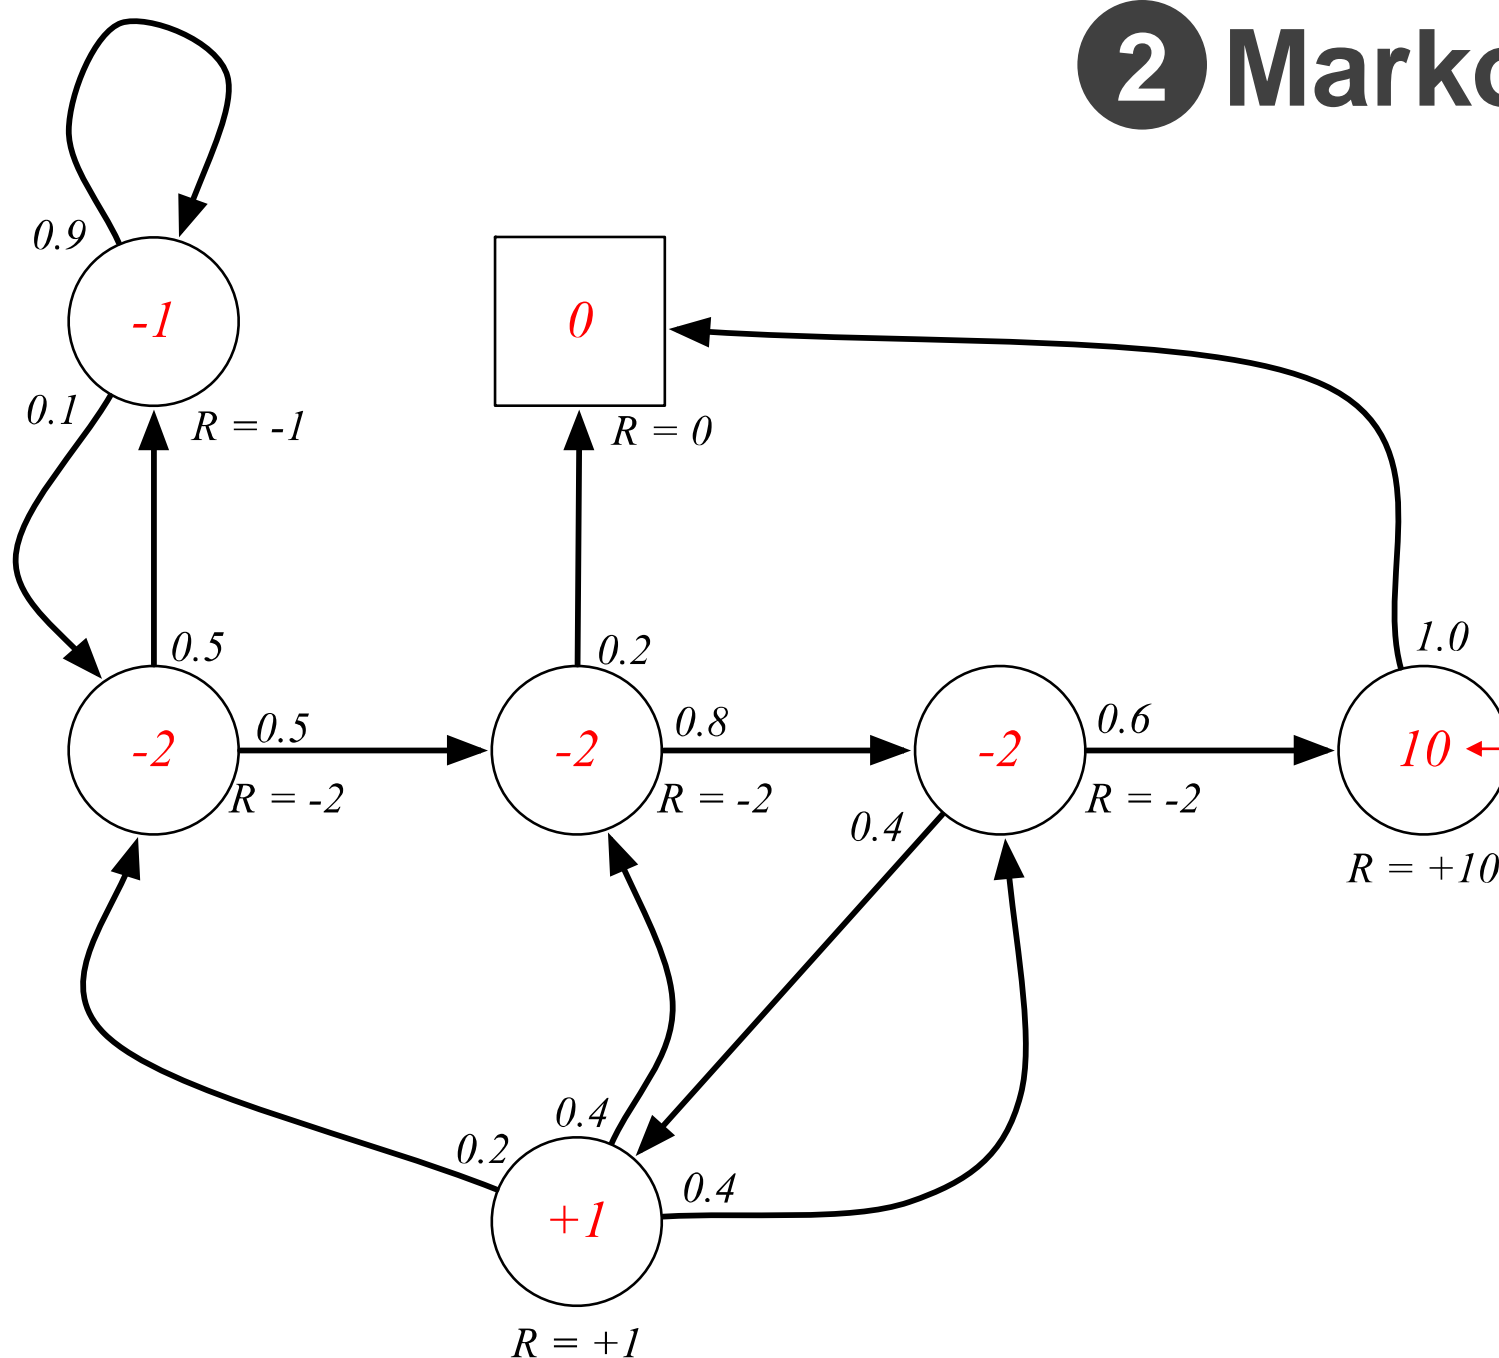
### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



$v(s)$  for  $\gamma = 0$

State value function  $v(s)$  is the expected total reward (into the future)

$$v(s) = E[G_t | S = s_t]$$

Example from David Silver, UCL, 2015

## 2 Markov Reward Process

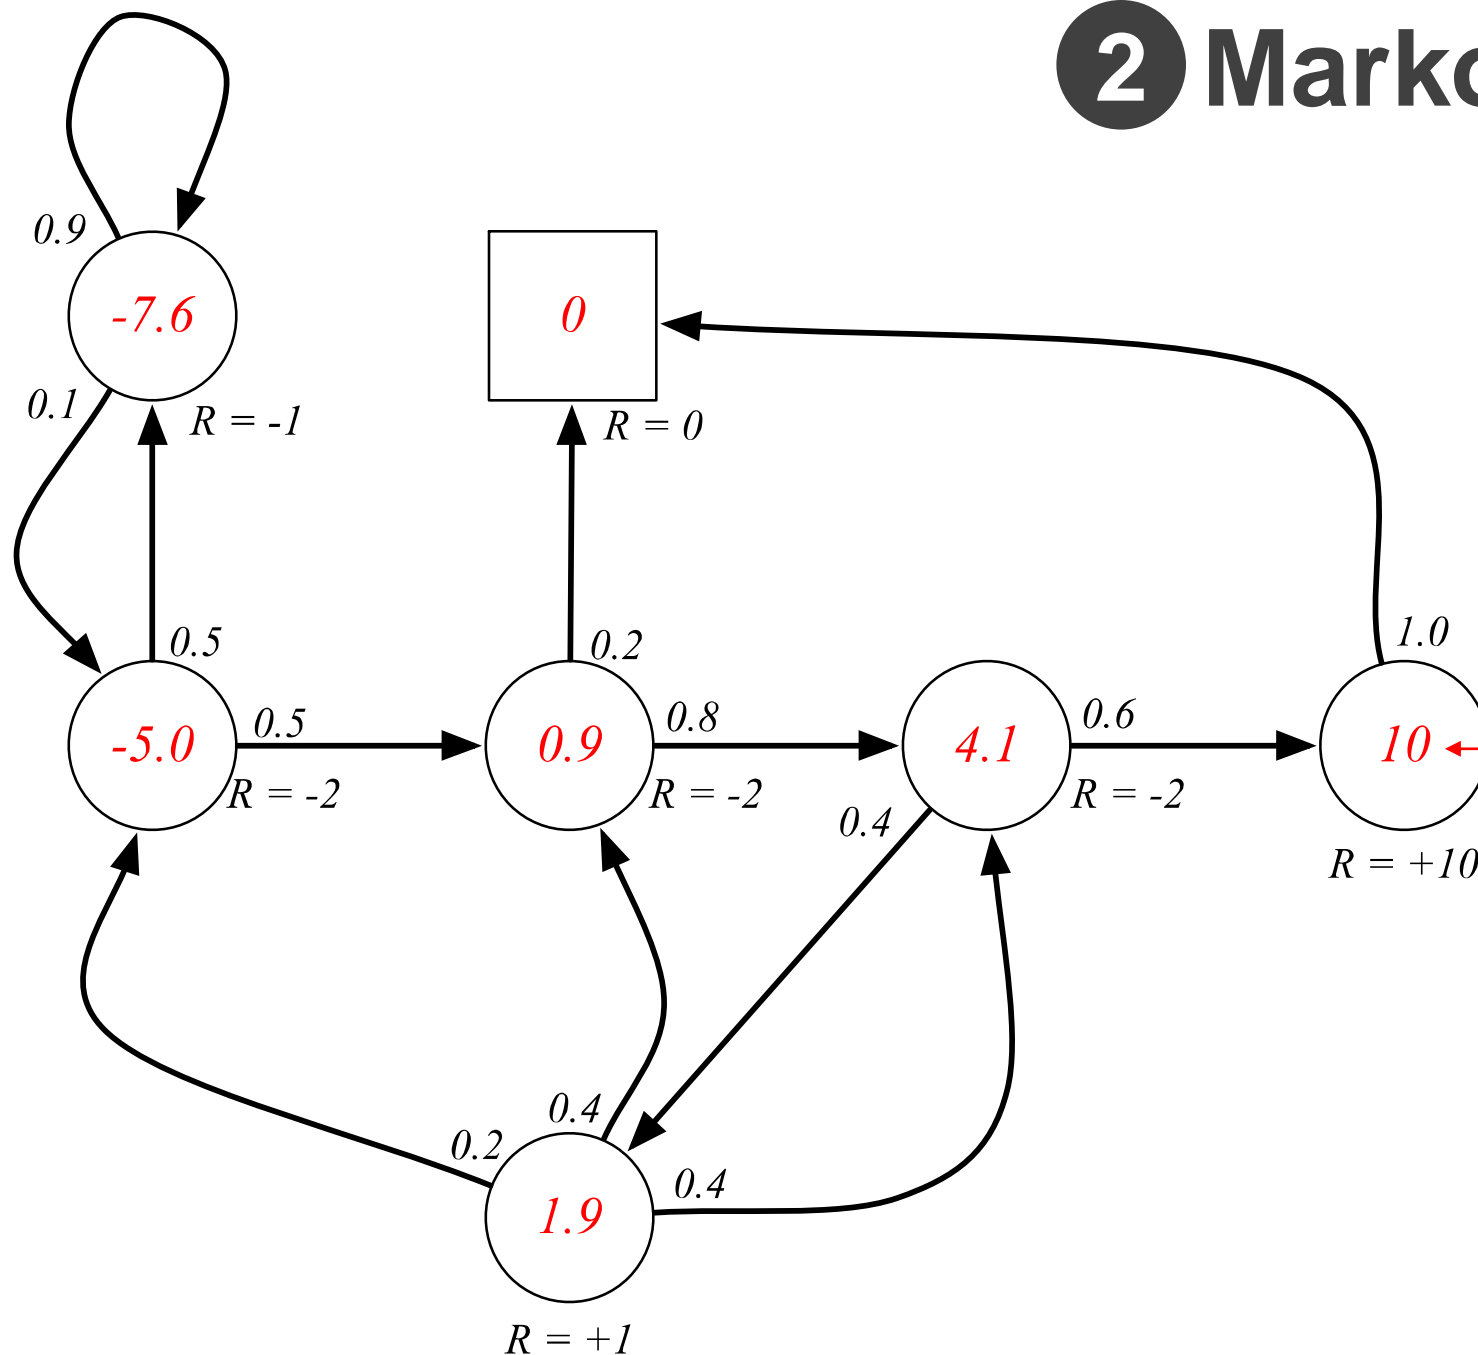
### Components:

State space  $S$ ,

Transition probabilities,  $P$

Rewards,  $R$

Discount rate,  $\gamma$



$v(s)$  for  $\gamma = 0.9$

State value function  $v(s)$  is the expected total reward (into the future)

$$v(s) = E[G_t | S = s_t]$$

Example from David Silver, UCL, 2015



# “Backup” property of state value functions

$$v(s) = E[G_t | S = s_t] \quad \text{where } G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots$$

$$= E[R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} \dots | S = s_t]$$

$$= E[R_{t+1} + \gamma(R_{t+2} + \gamma R_{t+3} \dots) | S = s_t]$$

$$= E[R_{t+1} + \gamma G_{t+1} | S = s_t]$$


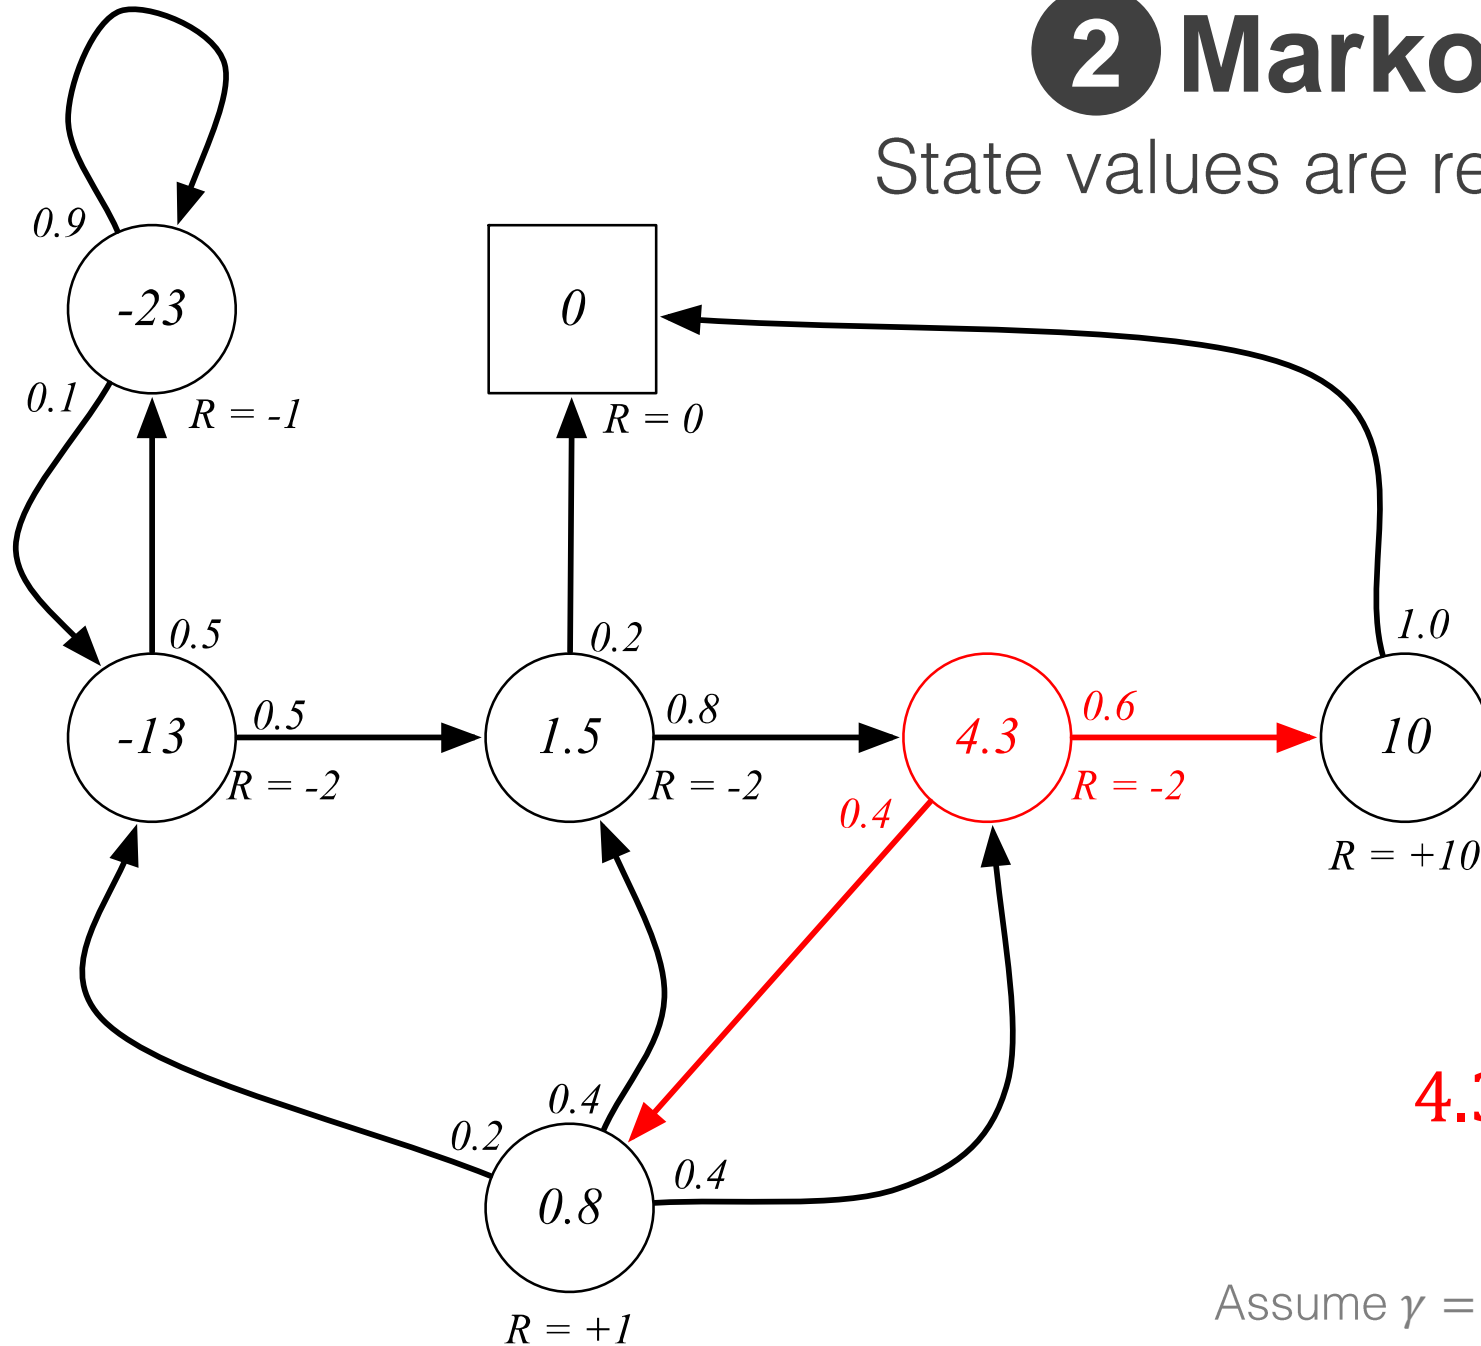
$$= E[R_{t+1} + \gamma v(s_{t+1}) | S = s_t]$$

This recursive relationship is a version of the **Bellman Equation**

Example from David Silver, UCL, 2015

## 2 Markov Reward Process

State values are related to neighboring states



possible states we could transition to from  $s$

$$v(s) = E[R_s + \gamma v(s') | s]$$

$$v(s) = R_s + \gamma \sum_{s'} P_{ss'} v(s')$$

$$4.3 = -2 + 0.6 \times 10 + 0.4 \times 0.8$$

Notation:  $s = s_t$  and  $s' = s_{t+1}$

$R_s = E[R_{t+1} | S_t = s]$

Assume  $\gamma = 1$

Example from David Silver, UCL, 2015

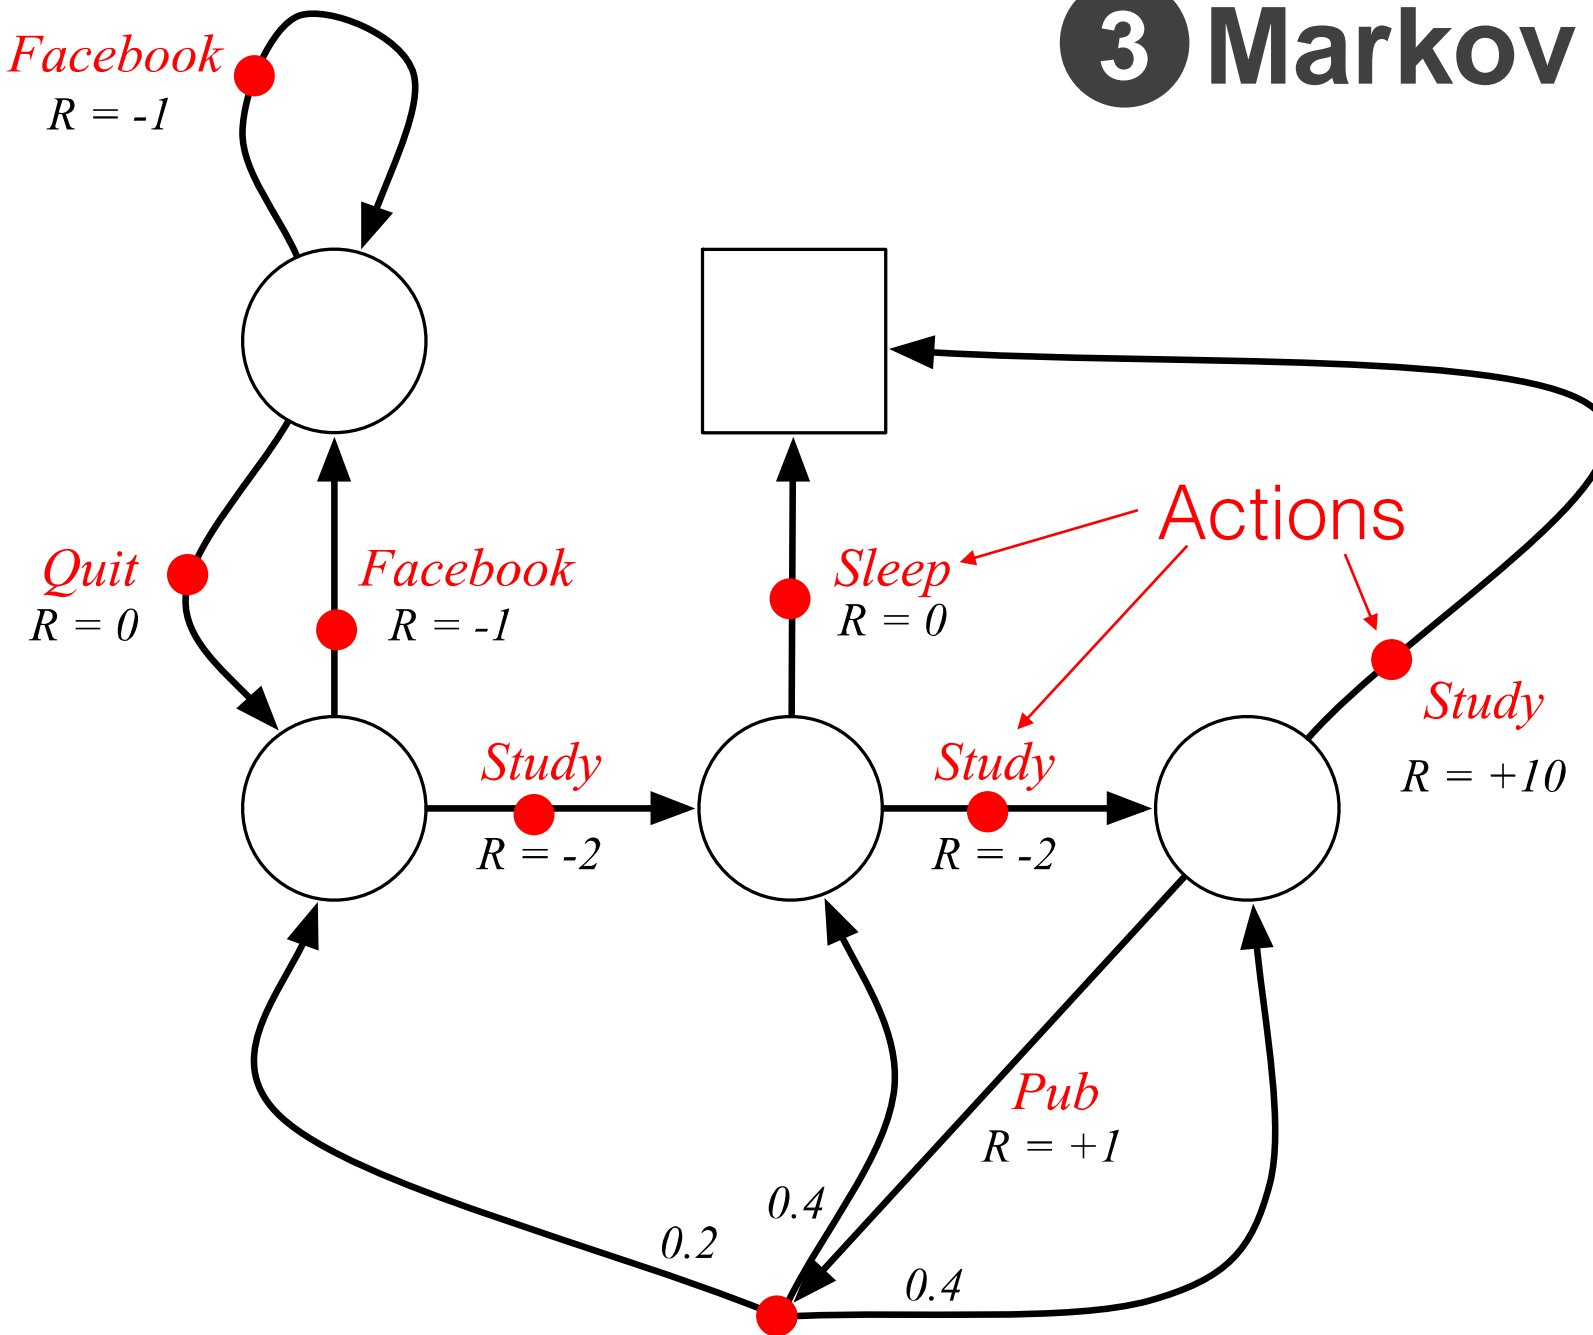
# 3 Markov Decision Process

## Components:

State space  $S$ ,  
Transition probabilities,  $P$   
Rewards,  $R$   
Discount rate,  $\gamma$   
Actions,  $A$

Adds interaction with the environment

An agent in a state chooses an action, the environment (the MDP) provides a reward and the next state



Example from David Silver, UCL, 2015

# 3 Markov Decision Process

Policy (how we choose actions)

(can be stochastic or deterministic)

$$\pi(a|s) = P(a|s)$$

State value function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E[G_t|s]$$

$$v_{\pi}(s) = E[R_s^a + \gamma v_{\pi}(s')|s]$$

Action value function

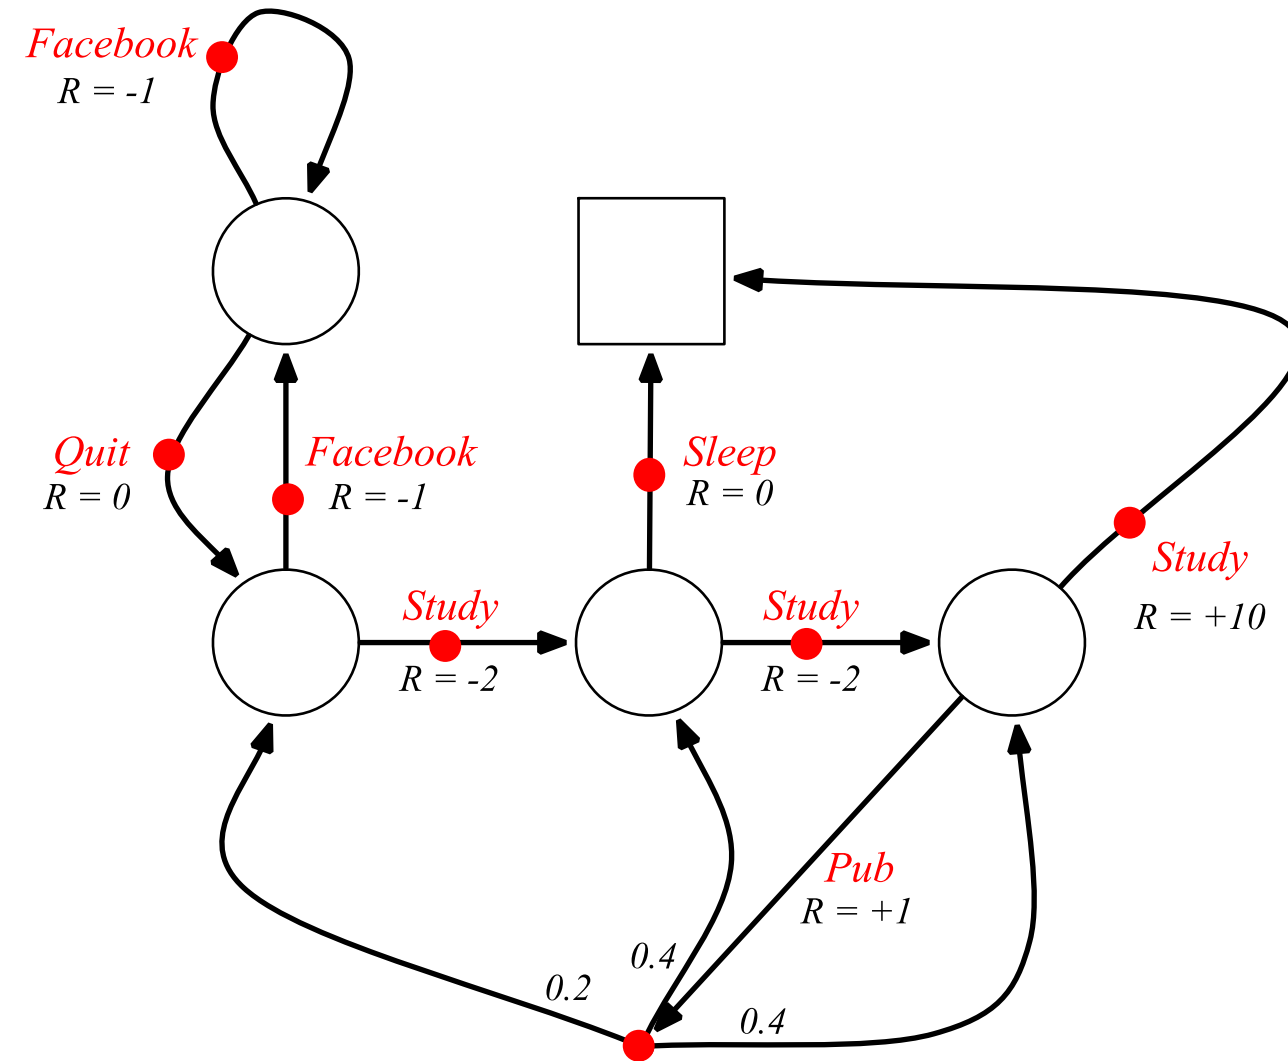
(expected return from state  $s$ , taking action  $a$ , and following policy  $\pi$ )

$$q_{\pi}(s, a) = E[G_t|s, a]$$

$$q_{\pi}(s, a) = E[R_s^a + \gamma q_{\pi}(s', a')|s, a]$$

$$R_s^a = E[r_{t+1}|S_t = s, A_t = a]$$

Example from David Silver, UCL, 2015




# Bellman Expectation Equations for the **state value** function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E[G_t | s]$$

$$v_{\pi}(s) = E[R_s^a + \gamma v_{\pi}(s') | s]$$

$$R_s^a = E[R_{t+1} | S_t = s, A_t = a]$$



Expectation over the possible actions

Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

$$v_{\pi}(s) = \underbrace{\sum_a}_{\text{green}} \underbrace{\pi(a|s)}_{\text{orange}} \left( \underbrace{R_s^a + \gamma \sum_{s'} P_{ss'}^a v_{\pi}(s')}_{\text{blue}} \right)$$

# Bellman Expectation Equations for the **action value** function


(expected return from state  $s$ , taking action  $a$ , then following policy  $\pi$ )

$$q_{\pi}(s, a) = E[G_t | s, a]$$

$$q_{\pi}(s, a) = E[R_s^a + \gamma q_{\pi}(s', a') | s, a]$$

$$R_s^a = E[R_{t+1} | S_t = s, A_t = a]$$

$$q_{\pi}(s, a) \leftarrow s, a$$



Expectation over the rewards

(based on state and choice of action)

Expectation over the next possible states

Expectation over the possible actions

$$q_{\pi}(s, a) = \underbrace{R_s^a}_{\text{orange}} + \gamma \underbrace{\sum_{s'} P_{ss'}^a}_{\text{blue}} \underbrace{\sum_{a'} \pi(a' | s') q_{\pi}(s', a')}_{\text{green}}$$


# Bellman Expectation Equations

## State value function

(expected return from state  $s$ , and following policy  $\pi$ )

$$v_{\pi}(s) = E[G_t | s]$$

$$v_{\pi}(s) = E[R_s^a + \gamma v_{\pi}(s') | s]$$




$$v_{\pi}(s) = \sum_a \pi(a|s) \left( R_s^a + \gamma \sum_{s'} P_{ss'}^a v_{\pi}(s') \right)$$

## Action value function

(expected return from state  $s$ , taking action  $a$ , then following policy  $\pi$ )

$$q_{\pi}(s, a) = E[G_t | s, a]$$

$$q_{\pi}(s, a) = E[R_s^a + \gamma q_{\pi}(s', a') | s, a]$$



$$q_{\pi}(s, a) = R_s^a + \gamma \sum_{s'} P_{ss'}^a \sum_{a'} \pi(a'|s') q_{\pi}(s', a')$$

# 3 Markov Decision Process

State value function,  $v_{\pi}(s)$

Assumptions:

$$\pi(a|s) = \frac{1}{\text{\# of reachable states}}$$

(randomly select an action)

$$\gamma = 1$$

(no discounting)


Bellman Expectation Equation

$$v(s) = E[R_s^a + \gamma v(s')|s]$$

$$-1.3 \approx 0.5 \times (-1 - 2.3) + 0.5 \times (-2 + 2.7)$$

Assume  $\gamma = 1$

Example from David Silver, UCL, 2015






# 3 Markov Decision Process

Optimal **state-value** function,  $v_*(s)$

Maximum value function over all policies

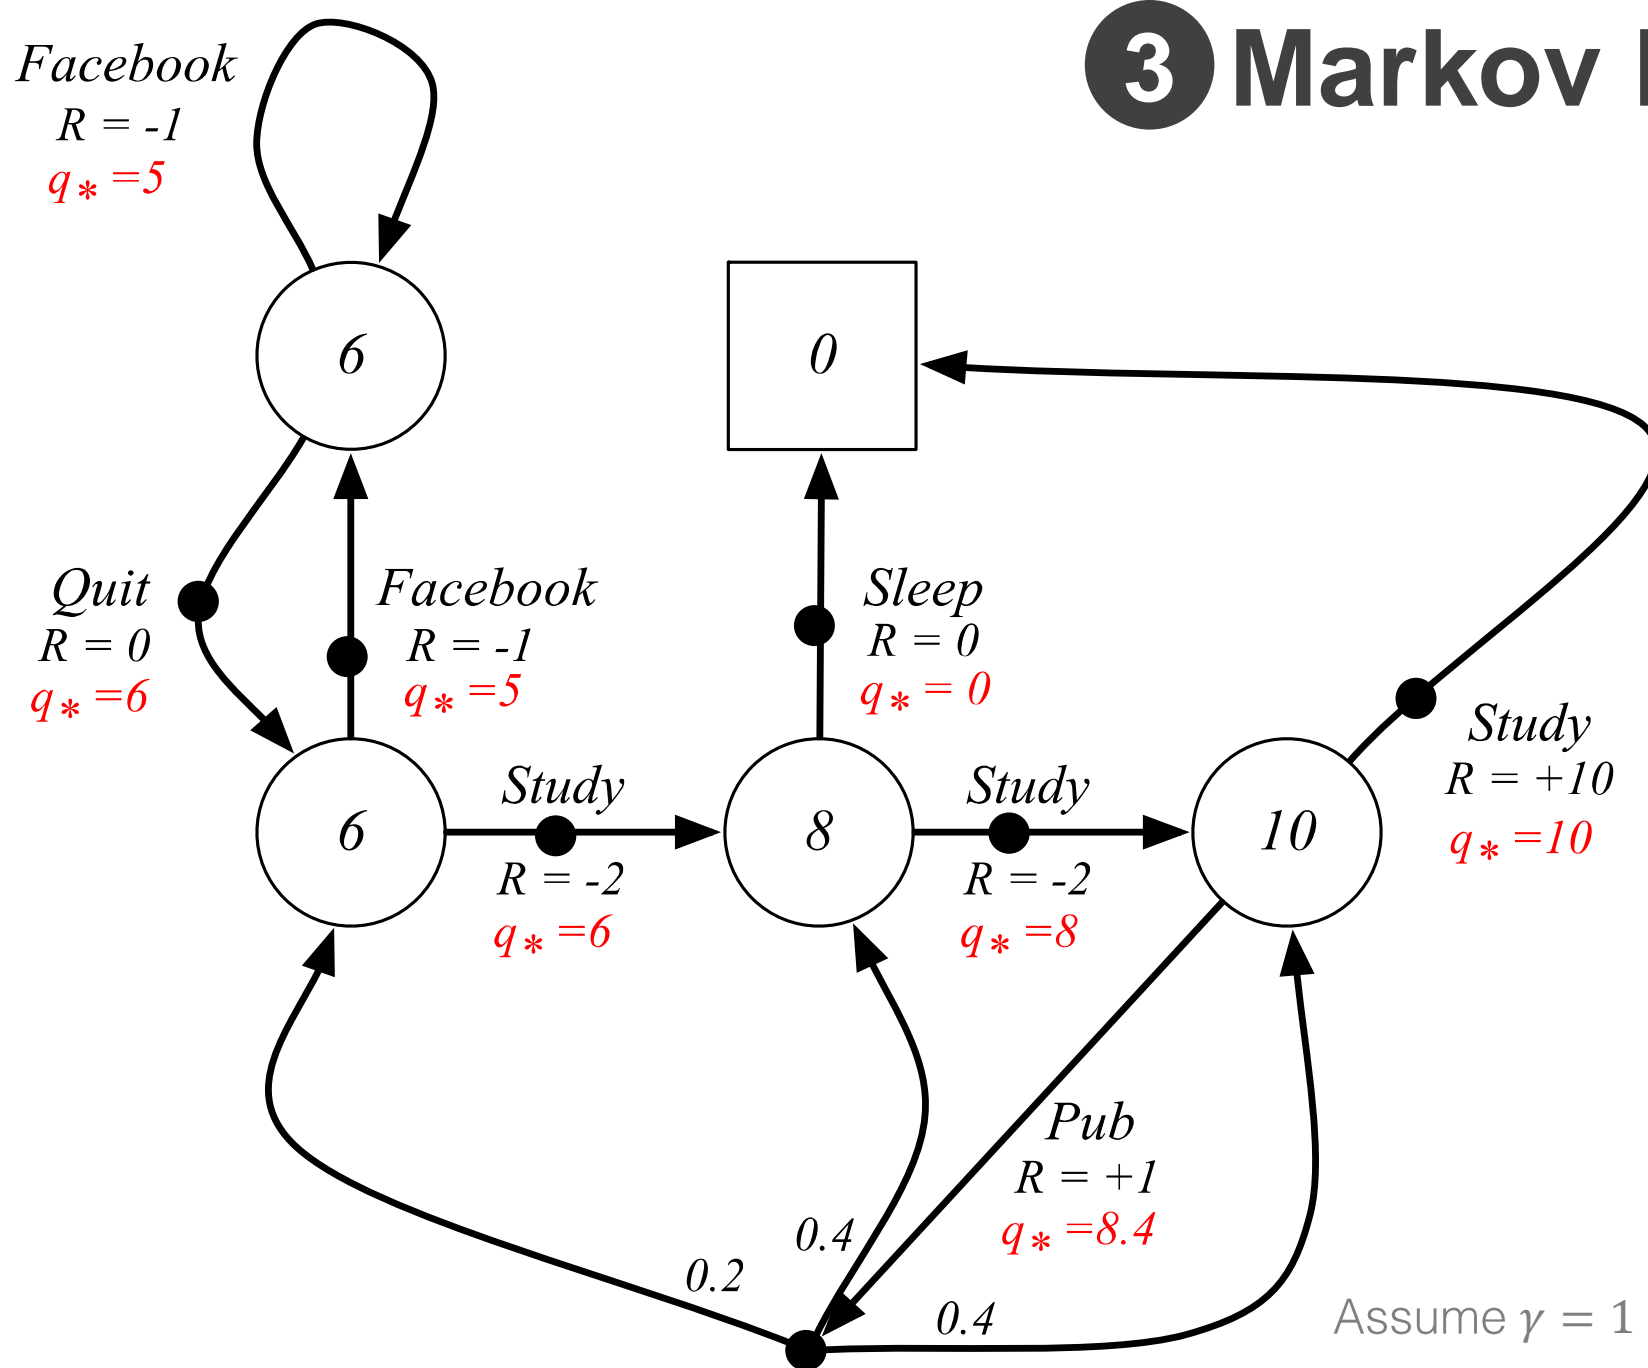
$$v_*(s) = \max_{\pi} v_{\pi}(s)$$



Assume  $\gamma = 1$

Example from David Silver, UCL, 2015

# 3 Markov Decision Process



Optimal **state-value** function,  $v_*(s)$

Maximum value function over all policies

$$v_*(s) = \max_{\pi} v_{\pi}(s)$$

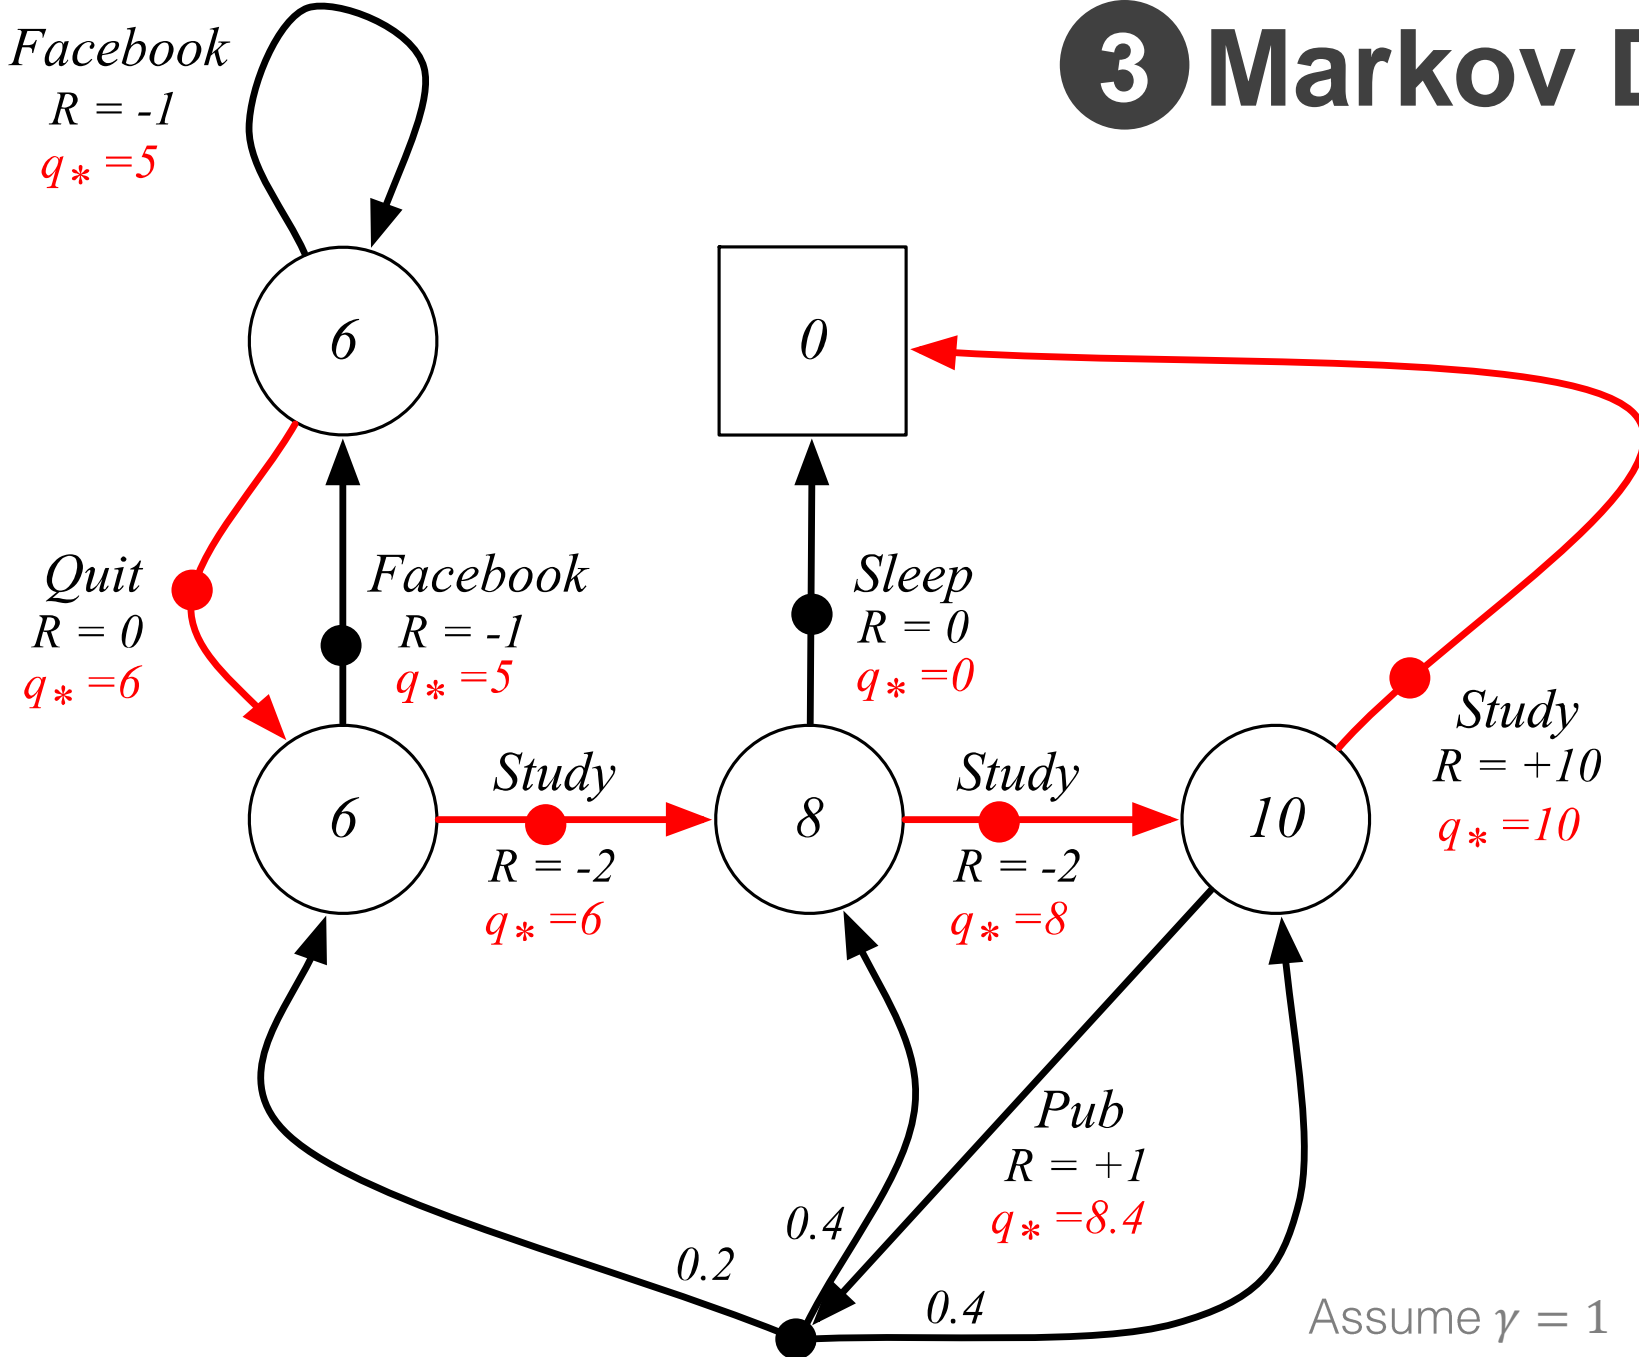
Optimal **action-value** function,  $q_*(s, a)$

Maximum value function over all policies

$$q_*(s, a) = \max_{\pi} q_{\pi}(s, a)$$

Example from David Silver, UCL, 2015

# 3 Markov Decision Process



Optimal **policy**,  $\pi_*(s)$   
Which action to take at each moment

$$\pi_*(s) = \arg \max_a q_*(s, a)$$

Once we have the optimal value functions, we've "solved" the MDP!

Example from David Silver, UCL, 2015

# Learning strategy

Model-based  
(planning)

Model-free

Reinforcement Learning

Knowledge of **Environment**

**No knowledge**  
Must learn from  
experience

**Monte Carlo Learning**

**Perfect  
knowledge**  
Known MDP

**Dynamic Programming**

Policy iteration  
Value iteration

**Next class:**

1. How to **compute optimal policies** for known MDPs?
2. How to extend this to the case **without full knowledge** of MDPs

